

General AEO, GEO, and SEO Client Playbook

Version: 1.0

Reviewed: 2026-07-23

Purpose

This is a reusable client playbook for modern search visibility across:

- Traditional Google and Bing search
- Google AI Overviews and AI search experiences
- ChatGPT, Perplexity, Gemini, Copilot, Claude, and similar assistants
- Agent-driven discovery flows where systems compare, summarize, and recommend before a user clicks

It is built from the recurring ideas shown in the supplied screenshots and shaped into a practical operating model for real client work.

This playbook does not promise rankings, citations, traffic, leads, sales, bookings, or inclusion in any AI answer. Its purpose is to improve the odds that a client can be:

- Found
- Fetched
- Understood
- Trusted
- Selected

The big shift

Traditional SEO was mostly about being found.

Modern AI-influenced search is about being chosen.

That means a client now has to win at five layers at once:

1. Crawlability
2. Page clarity
3. Structured data
4. Trust signals
5. Conversion readiness

If a system cannot clearly understand what the client offers, why it matters, and whether it is safe to recommend, the client will struggle even if some pages technically rank.

The simple model: SEO vs AEO vs GEO

Use these terms in a practical way:

- SEO = help search engines find and rank the page
- AEO = help answer engines understand and extract the page
- GEO = help generative systems trust, justify, and cite the page

Another way to think about it:

- SEO drives discovery
- AEO drives understanding
- GEO drives confidence

You need all three.

How AI-influenced discovery works

The screenshots repeatedly point to a seven-stage pattern. Treat this as a working model for planning content.

Stage 1: Query classification

The system first decides whether the query needs:

- memory only
- simple search
- complex multi-step search

High-intent queries usually trigger search, especially queries about:

- choosing a provider
- comparing options
- evaluating alternatives
- making a purchase
- booking a service

If the client is not meaningfully present in search, they may never enter the AI pipeline.

Stage 2: Search planning

The system then expands the query into related searches, often including:

- broad keyword-style searches

- intent-weighted long-tail searches
- comparison searches
- “best for” searches
- location and use-case variants

This means pages should not only target one broad keyword. They must cover the full decision angle around the query.

Stage 3: Candidate page selection

The system pulls many possible pages, then filters aggressively using signals such as:

- title clarity
- description clarity
- apparent relevance
- trust signals
- diversity

Most pages get discarded here.

Stage 4: Fetching

If the page is slow, unstable, blocked, or difficult to render, it may be:

- truncated
- only partially read
- dropped entirely

Speed is no longer just a UX issue. It affects whether the content reaches the system at all.

Stage 5: Chunking and semantic scoring

The system does not “read the page like a person.” It breaks the page into chunks and scores those chunks against intent.

Only the strongest chunks survive.

That is why formatting, headings, answer blocks, tables, and Q&A structure matter so much.

Stage 6: Final page selection

The system narrows to a small candidate set and may only deeply read a handful of pages.

If your page does not survive early filtering, it does not matter how good the later sections are.

Stage 7: Final synthesis

The final answer may blend:

- extracted page chunks
- trusted summaries
- reviews
- third-party mentions
- news
- maps
- live pricing or availability

If you fail the earlier stages, you do not appear in the answer.

The three data layers every client must control

The screenshots repeatedly describe three distinct data planes. This is one of the most useful frameworks in the whole set.

1. Crawled web data

This includes:

- what search engines index
- what AI systems can find on the public web
- what they learned from prior crawling and public content

This shapes baseline brand perception.

SEO still matters here.

2. Structured feeds and machine-readable data

This includes:

- schema markup
- product feeds
- inventory feeds
- price and availability fields
- category and entity relationships
- APIs where relevant

This is where precision and control live.

For ecommerce especially, this often drives:

- comparisons
- rankings
- recommendations

3. Live website data

This includes what a system sees when it actually visits the page, such as:

- current offer details
- pricing
- promotions
- reviews
- media
- forms
- booking UX
- checkout capability

If the live site fails, the recommendation path often fails with it.

The three operating pillars

Pillar 1: Technical foundations and structured data

AI systems need clean, consistent, machine-readable structure.

Core requirements:

- semantic HTML
- correct canonicals
- valid metadata
- correct indexing rules
- fast rendering
- accessible components
- stable layout
- structured data that matches visible content

For many clients, this is where the biggest opportunity is because the data exists but is not exposed clearly enough.

Pillar 2: Intent-driven content enrichment

AI increasingly interprets intent, not just exact keywords.

Pages should front-load:

- who the offer is for
- what problem it solves
- why it is better
- when to choose it
- when not to choose it

Use:

- question-based H2s
- direct answers early in sections
- FAQ sections
- comparison blocks
- feature lists
- use-case framing
- “goes well with” relationships where relevant
- video transcripts where relevant
- detailed image alt text

Pillar 3: Trust and credibility signals

A system must feel safe recommending the client.

Trust comes from:

- verified reviews
- review volume
- transparent policies
- expert authorship
- press mentions
- certifications
- official profiles
- strong entity consistency
- accurate claims
- up-to-date content

Do not exaggerate. AI systems are increasingly sensitive to unsupported certainty.

What this means for client strategy

The winning clients treat:

- data as a product
- feeds as strategic assets
- content as machine-readable infrastructure
- trust as a measurable ranking factor

The old mindset was “rank and hope.”

The new mindset is “be the obvious, safe, well-structured choice.”

Universal page model for clients

Every important service, offer, course, product, or location page should answer these questions near the top:

1. What is this?
2. Who is it for?
3. What problem does it solve?
4. Why choose this over alternatives?
5. What is included?
6. What does it cost or how is pricing handled?
7. What should the visitor do next?

Recommended page structure:

1. H1 with the actual offer or question
2. Direct answer paragraph
3. Who it is for
4. Problem solved
5. Why it is different
6. Inclusions and exclusions
7. Pricing or pricing framework
8. Process or how it works
9. Proof and trust signals
10. FAQ
11. CTA

Question-based H2 framework

The screenshots strongly support question-led section design.

Good H2 patterns:

- What does this service include?

- Who is this best for?
- How much does this cost?
- How long does the process take?
- Why choose this over other options?
- What results should a client realistically expect?
- What happens after I book?
- Is this right for beginners?
- What is the difference between option A and option B?

These work well because they match real search and AI prompt behavior.

Content blocks AI can extract easily

Direct answer block

Start a section with a plain-language answer in 40 to 80 words where natural.

Comparison table

Use when a query implies evaluation.

TL;DR summary

Use near the top for higher-friction topics.

FAQ block

Use visible FAQs with plain answers. Then add FAQ schema only if the FAQs are truly visible on the page.

Feature list

Use bullets for scannable extraction.

Proof block

Put evidence close to the claim it supports.

Internal relationship block

Use sections like:

- Best for
- Works well with

- Compare with
- Related services
- Related courses

These help systems understand entity relationships.

Content depth vs content scale

One well-written page is not enough anymore for many competitive markets.

The screenshots repeatedly point to topical coverage and fan-out query coverage.

A client needs:

- a strong primary service page
- supporting comparison pages
- FAQs
- use-case pages
- articles answering adjacent questions
- proof pages or case studies

Coverage often beats isolated rankings because it embeds the client more deeply into the topic.

Schema strategy by client type

Use only schema that matches visible reality.

All clients

Use:

- `Organization`
- `WebSite`
- `BreadcrumbList`
- `FAQPage` when applicable
- `Article` or `BlogPosting` for editorial content

Service businesses

Use:

- `Service`
- `LocalBusiness` if location-based
- `ProfessionalService` where appropriate

- `FAQPage`

Education businesses

Use:

- `EducationalOrganization`
- `Course`
- `CourseInstance` where applicable
- `FAQPage`
- `Person` for instructors

Ecommerce businesses

Use:

- `Product`
- `Offer`
- `AggregateRating`
- `Review`
- `Brand`
- `ItemList` for category pages
- `ImageObject` where useful

Important feed-like fields called out repeatedly in the screenshots:

- price
- availability
- size
- color
- SKU
- GTIN
- dateModified
- priceCurrency
- inLanguage

Media-heavy pages

Use:

- `VideoObject`
- `ImageObject`

Minimum technical baseline

Every client should have:

- crawlable public pages
- clean robots.txt
- XML sitemap
- canonical URLs
- no accidental noindex
- no cloaking
- no bot-only HTML
- fast server response
- mobile usability
- stable layout
- accessible forms and buttons
- indexable primary copy in HTML
- clear internal linking

Trust stack checklist

Every serious client should strengthen as many of these as truthfully possible:

- real reviews
- review schema when eligible
- clear contact details
- visible policies
- author names and credentials
- expert bylines
- citations to primary sources
- brand mentions on trusted third-party sites
- certifications or memberships
- press mentions
- consistent brand naming everywhere
- updated content dates

Reusable content brief for any client

Before writing, define:

- Primary decision query
- Secondary supporting queries
- Audience
- Geography
- Intent type: informational, comparison, transactional, local, educational
- Desired conversion

- Main proof available
- Main objections
- Main differentiators
- Which schema types apply

Fast-start client roadmap

Week 1

- Audit technical SEO baseline
- Identify top 5 money pages
- Identify missing schema
- Identify top 10 client questions
- Review crawl/indexation issues

Week 2

- Rewrite top money pages using direct-answer structure
- Add question-led H2s
- Add proof sections
- Add FAQs
- Fix metadata and canonicals

Week 3

- Add or repair schema
- Improve internal linking
- Add comparison and “best for” support content
- Improve page speed and rendering stability

Week 4

- Publish supporting cluster content
- Add trust assets
- Review AI and search visibility manually
- Measure conversions, not just rankings

What not to do

Do not:

- promise rankings
- promise ChatGPT inclusion
- publish fake reviews

- invent aggregate ratings
- use schema not supported by visible content
- create near-duplicate keyword pages
- hide content from users and show it to bots
- stuff keywords unnaturally
- make exaggerated “best in the world” claims without evidence
- use stale pricing in copy or schema
- rely on one thin page in a competitive market

The reusable client standard

For any client, the goal is not “get indexed and hope.”

The goal is:

- Make the client easy to understand
- Make the client easy to compare
- Make the client easy to trust
- Make the client easy to select
- Make the next step easy to complete

That is the real overlap between SEO, AEO, and GEO.

Final takeaway

If a client cannot be clearly interpreted, justified, and acted on in real time, they are weak in modern search.

The best future-proof strategy is:

- strong technical SEO
- question-led answer-first content
- correct schema
- deep topical coverage
- real trust signals
- frictionless conversion paths

That combination gives clients the best chance to show up across both traditional search and AI-mediated discovery.